# **Day 5 – Kubernetes Architecture**

### 1. Explain Kubernetes Architecture in depth.

Ans) Kubernetes has a modular and distributed architecture, designed for scalability, resilience, and declarative automation. At a high level, Kubernetes is made up of a control plane (master components) and a data plane (worker nodes).

# 🔆 Components in Detail

A. Control Plane (Master Components)

#### a. kube-apiserver

- Acts as the gateway to the Kubernetes cluster.
- All interactions (via kubectl, other tools, or internal components) go through this.
- Exposes the Kubernetes API (REST) for clients/developers to interact.

#### b. etcd

- A distributed, consistent key-value store.
- Stores all cluster data: nodes, pods, secrets, config maps, etc.
- Highly available and fault-tolerant.

#### c. kube-scheduler

- Watches for newly created Pods without assigned nodes.
- Selects the best node to run a Pod based on Resource Availability, Affinity Rules, Taints and Tolerations, Node Selectors etc.

## d. kube-controller-manager

- Runs controller loops to reconcile cluster state.
- Includes:
- i) Node Controller detects and reacts to node failures.
- ii) Replication Controller maintains desired number of pod replicas.
- iii) Job Controller, Endpoint Controller, etc.
- B. Node Components (Worker Nodes)

## a. kubelet

- Agent that runs on each worker node.
- Ensures containers are running in a Pod.
- Communicates with the control plane.

# b. kube-proxy

- Handles networking and traffic routing.
- Maintains network rules on nodes.
- Enables service discovery and load balancing.

## c. Container Runtime

- The software that actually runs containers (e.g., containerd, CRI-O, or Docker)
- Interfaces with kubelet via the Container Runtime Interface (CRI).

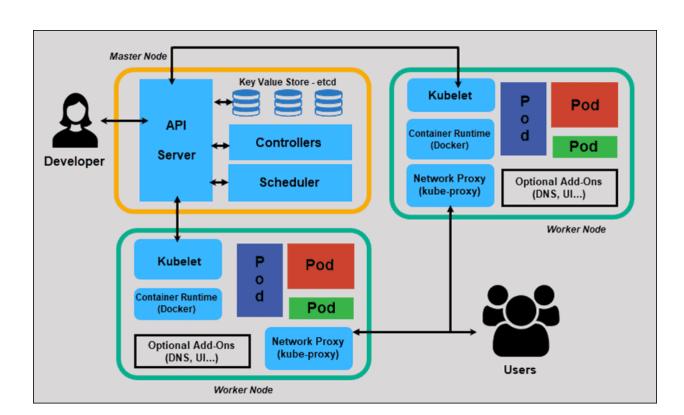
## d. Pod:

The smallest deployable unit in Kubernetes.

if they fail or are evicted.

- A Pod can hold one or more containers that share network namespace (same IP/ port space), Storage Volumes, Lifecycle (Start/Stop together)
- Pods are ephemeral; Kubernetes replaces them automatically

- How It All Works Together (Flow)
- a. User submits deployment via kubectl (talks to kube-apiserver).
- b. kube-apiserver validates and stores the object in etcd.
- c. Controllers notice the desired state and take action (e.g., create Pods).
- d. Scheduler assigns Pods to appropriate nodes.
- e. kubelet on each node receives the Pod spec and launches containers.
- f. kube-proxy sets up routing so services can reach the right Pods.



### 2. What are the different functionalities of kube-apiserver?

Ans) The kube-apiserver is the central management component and entry point for all operations in a Kubernetes cluster. It exposes the Kubernetes API, which is used by all components—internal (like controllers, schedulers) and external (like kubectl, CI/CD tools)—to communicate with and manage the cluster.

- Core Functionalities of kube-apiserver
- a. 📝 Authentication
- -> Validates the identity of users and components making requests.
- -> Supports multiple authentication methods:
  - Certificates
  - Bearer Tokens
  - Service Accounts
  - OIDC (OAuth2)
  - Webhooks
- b. 🗸 Authorization
- -> Once authenticated, determines what actions the entity is allowed to perform.
- -> Uses:
  - Role-based Access Control (RBAC)
  - Attribute-based Access Control (ABAC)
  - Webhook Authorization
  - Node Authorization
- c. 🔍 Admission Control
- -> Intercepts and validates or mutates requests after authentication
- and authorization.

-> Types of admission controllers:

- ValidatingAdmissionWebhook Validate the request.
- MutatingAdmissionWebhook Modify the object before saving.
- Examples: NamespaceLifecycle, LimitRanger, ResourceQuota
- d. API Handling & Validation
- -> Parses incoming requests (usually via REST).
- -> Validates API objects against schemas before storing.
- -> Converts between different API versions (e.g., apps/v1, v1beta1).
- e. State Persistence (etcd interaction)
- -> Stores and retrieves the desired state of the cluster from etcd.
- -> Every object created/modified through the API (Pods, Deployments, Services) is stored here.
- -> Acts as a source of truth for the cluster state.
- f. 🙎 Cluster Communication Hub
- -> Serves as the hub for internal component communication:
  - The scheduler fetches unscheduled Pods.
  - The controller manager monitors objects and ensures desired state.
  - kubelet watches for Pods assigned to its node.
- g. 👲 RESTful Interface for Clients
- -> Exposes REST APIs used by: kubectl CLI, Dashboard UI, CI/CD tools like ArgoCD, Gitlab CI, third party apps integrating with Kubernetes.

The kube-apiserver is like the brainstem of Kubernetes—every single cluster interaction passes through it.